

Below is code with a link to a happy or sad dataset which contains 80 images, 40 happy and 40 sad. Create a convolutional neural network that trains to 100% accuracy on these images, which cancels training upon hitting training accuracy of $>.999$

Hint -- it will work best with 3 convolutional layers.

In [2]:

```
import tensorflow as tf
import os
import zipfile
from os import path, getcwd, chdir

# DO NOT CHANGE THE LINE BELOW. If you are developing in a local
# environment, then grab happy-or-sad.zip from the Coursera Jupyter Notebook
# and place it inside a local folder and edit the path to that location
path = f"{getcwd()}/../tmp2/happy-or-sad.zip"

zip_ref = zipfile.ZipFile(path, 'r')
zip_ref.extractall("/tmp/h-or-s")
zip_ref.close()
```

In [4]:

```
# GRADED FUNCTION: train_happy_sad_model
def train_happy_sad_model():
    # Please write your code only where you are indicated.
    # please do not remove # model fitting inline comments.

    DESIRED_ACCURACY = 0.999

    class myCallback(tf.keras.callbacks.Callback):
        def on_epoch_end(self, epoch, logs={}):
            if(logs.get('acc')>DESIRED_ACCURACY):
                print("\nReached 99.9% accuracy so cancelling training!")
                self.model.stop_training = True

    callbacks = myCallback()

    # This Code Block should Define and Compile the Model. Please assume the images
    # are 150 X 150 in your implementation.
    model = tf.keras.models.Sequential([
        # Your Code Here
        tf.keras.layers.Conv2D(16, (3,3), activation='relu', input_shape=(150, 150,
3)),
        tf.keras.layers.MaxPooling2D(2, 2),
        tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
        tf.keras.layers.MaxPooling2D(2,2),
        tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
        tf.keras.layers.MaxPooling2D(2,2),
        tf.keras.layers.Flatten(),
        tf.keras.layers.Dense(512, activation='relu'),
        tf.keras.layers.Dense(1, activation='sigmoid')

    ])

    from tensorflow.keras.optimizers import RMSprop

    model.compile(loss = 'binary_crossentropy',
                  optimizer = RMSprop(lr=0.001),
                  metrics = ['accuracy'])

    # This code block should create an instance of an ImageDataGenerator called
    # train_datagen
    # And a train_generator by calling train_datagen.flow_from_directory

    from tensorflow.keras.preprocessing.image import ImageDataGenerator

    train_datagen = ImageDataGenerator(rescale=1/255)

    # Please use a target_size of 150 X 150.
    train_generator = train_datagen.flow_from_directory(
        "/tmp/h-or-s",
        target_size=(150, 150),
        batch_size=10,
        class_mode='binary')

    # Expected output: 'Found 80 images belonging to 2 classes'

    # This code block should call model.fit_generator and train for
    # a number of epochs.
    # model fitting
```

```
history = model.fit_generator(  
    train_generator,  
    steps_per_epoch=8,  
    epochs=15,  
    verbose=1,  
    callbacks=[callbacks])  
  
# model fitting  
return history.history['acc'][-1]
```

In [5]:

```
# The Expected output: "Reached 99.9% accuracy so cancelling training!"  
train_happy_sad_model()
```

WARNING: Logging before flag parsing goes to stderr.
W0422 07:14:13.570796 140280201516864 deprecation.py:506] From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/init_ops.py:1251: calling VarianceScaling.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

W0422 07:14:13.967475 140280201516864 deprecation.py:323] From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/nn_impl.py:180: add_dispatch_support.<locals>.wrapper (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Found 80 images belonging to 2 classes.

Epoch 1/15

8/8 [=====] - 6s 699ms/step - loss: 1.0374 - acc: 0.6000

Epoch 2/15

8/8 [=====] - 0s 62ms/step - loss: 0.6159 - acc: 0.6750

Epoch 3/15

8/8 [=====] - 0s 62ms/step - loss: 0.2063 - acc: 0.9375

Epoch 4/15

8/8 [=====] - 1s 63ms/step - loss: 0.2023 - acc: 0.9125

Epoch 5/15

8/8 [=====] - 0s 62ms/step - loss: 0.3431 - acc: 0.8875

Epoch 6/15

8/8 [=====] - 1s 63ms/step - loss: 0.0859 - acc: 0.9750

Epoch 7/15

8/8 [=====] - 0s 62ms/step - loss: 0.0636 - acc: 0.9750

Epoch 8/15

8/8 [=====] - 0s 53ms/step - loss: 0.2807 - acc: 0.9125

Epoch 9/15

8/8 [=====] - 0s 54ms/step - loss: 0.0661 - acc: 0.9750

Epoch 10/15

8/8 [=====] - 0s 53ms/step - loss: 0.0291 - acc: 0.9875

Epoch 11/15

8/8 [=====] - 0s 53ms/step - loss: 0.0617 - acc: 0.9875

Epoch 12/15

8/8 [=====] - 0s 53ms/step - loss: 0.0370 - acc: 0.9875

Epoch 13/15

7/8 [=====>....] - ETA: 0s - loss: 0.0105 - acc: 1.0000

Reached 99.9% accuracy so cancelling training!

8/8 [=====] - 0s 53ms/step - loss: 0.0096 - acc: 1.0000

Out[5]:

1.0

In [4]:

```
# Now click the 'Submit Assignment' button above.  
# Once that is complete, please run the following two cells to save your work and close the notebook
```

In []:

In []: